

NIKUNJ

+91 9996739385 | nikunjsachdeva919@gmail.com | linkedin.com/in/nikunjsachdeva | github.com/NikunjSachdeva | nikunj-sachdeva.vercel.app

EDUCATION

Manipal University Jaipur — Jaipur, Rajasthan B.Tech, Computer Science and Engineering (AI/ML) | CGPA: 8.12 | Graduation: 2026

TECHNICAL SKILLS

Languages: Python, C, C++, SQL, JavaScript, Solidity

Platforms: Windows Application Development, Linux

Frameworks & Libraries: FastAPI, Flask, React.js, LangChain, OpenAI SDK, Jinja2, Gradio, PyYAML

AI / ML: Machine Learning, Deep Learning, NLP, Large Language Models (LLMs), Generative AI, RAG, Computer Vision

Data & Databases: NumPy, Pandas, SQLite, MySQL, Pinecone, Firebase Firestore

Cloud & DevOps: GCP, AWS, Git, GitHub, REST APIs, CI/CD, Automation

Tools: VS Code, Confluence, Cppcheck, Chart.js, Cloudinary, OpenAPI

EXPERIENCE

Siemens Technology and Services Private Limited (STSPL) — Bangalore, India

Technical Intern | Feb 2026 – Jul 2026

- Built **CodePulse**, a Python CLI automation bot monitoring VCS for source code changes and auto-generating structured wiki documentation via an OpenAI-compatible LLM gateway from C++ headers and READMEs.
- Implemented **SHA-based incremental caching** that skips LLM inference for unchanged files, compressing a full 17-page regeneration (4 min 11 sec) down to ~42 sec for a single-driver change and 18 sec for no-change runs.
- Engineered an **HTML sanitizer pipeline** eliminating 100% of XHTML validation failures across 47 Confluence page push operations, ensuring enterprise-grade knowledge base integrity.
- Developed a **PR analytics pipeline** with SQLite and Jinja2 + Chart.js dashboards, delivering engineering quality insights to 100+ engineers.
- Integrated **Cppcheck static analysis** for complexity metrics, duplication detection, and CWE Top 25 SAST scanning; reduced documentation turnaround from 45–90 min/page to ~30 sec.

PROJECTS

CodePulse — Knowledge-Based Automated Documentation Bot | Python, OpenAI SDK, SQLite, Jinja2, Confluence API

- LLM-powered CLI tool that automatically generates and publishes XHTML-compliant wiki pages to Confluence from C++ source changes, eliminating manual documentation bottlenecks for 100+ engineers.
- Achieved 100% XHTML validation success rate and 98% reduction in documentation turnaround time through incremental caching and automated HTML sanitization.

Mini-RAG Pro — High-Performance RAG Backend | Python, FastAPI, Pinecone, Google Gemini, Cohere Rerank, Gradio

github.com/NikunjSachdeva/Mini-rag

- Built a production-grade Retrieval-Augmented Generation backend with semantic chunking, Pinecone vector search, and Cohere Rerank multi-stage relevance optimization, reducing query latency by 40%.
- Integrated Google Gemini LLM for citation-aware response generation with a Gradio UI, supporting scalable document ingestion pipelines.

DevTok — Full-Stack Developer Social Platform | FastAPI, React.js, Firebase Firestore, Cloudinary, LLMs

github.com/NikunjSachdeva/DevTok

- Architected a developer-focused platform with LLM-powered auto-generated quizzes and code explanations, real-time media feeds via Cloudinary, and Firestore-backed authentication.
- Built RESTful API with FastAPI and a React.js frontend, supporting multi-media content creation, user engagement features, and scalable cloud storage.

Advanced Rainfall Prediction — Deep Learning Forecasting System | Python, Deep Learning, NumPy, Pandas

github.com/NikunjSachdeva/rainfall_prediction

- Engineered a deep learning pipeline trained on 33 years of IMD meteorological data achieving MAE < 0.005 and accurately forecasting daily rainfall up to 77.1 mm/day.
- Applied feature engineering, time-series modeling, and rigorous evaluation on large-scale geospatial datasets for high-accuracy weather prediction.